

PAPER

Assessing the Reliability of AlphaFold3 Predictions for Protein-Ligand Affinity Prediction via Sfcnn

Guo Yu,¹ Yiming Wu¹ and Yiyang Tan¹¹School of Information Science and Technology, ShanghaiTech University, 393 Middle Huaxia Road, 201210, Shanghai, China

Abstract

This study evaluates the reliability of AlphaFold3 (AF3)-predicted protein structures for protein-ligand affinity (PLA) prediction tasks. We reproduce the Sfcnn model, a 3D convolutional neural network (CNN) for PLA prediction, using PyTorch. Model performance is validated on the PDBbind v2019 refined set for training and the CASF-2016 core set for testing. Subsequently, AF3-derived protein structures from the CASF-2016 core set are assessed and compared to ground truth structures using Sfcnn scores, to determine the suitability of AF3 predictions in PLA applications.

Key words: AlphaFold3, protein-ligand affinity, CNN scoring function, CASF-2016

Introduction

Sfcnn Overview

Sfcnn is a 3D convolutional neural network-based scoring function introduced in 2022 by Wang et al¹, designed to provide accurate and reliable predictions of binding affinities for protein-ligand complexes.

Methods

Datasets

The Sfcnn network was trained using protein-ligand complexes from the PDBbind v2019 refined set, which includes experimentally determined binding affinities (pKa values). The model was evaluated on the CASF-2016 core set, comprising 285 protein-ligand complexes. To prevent data leakage, 266 overlapping protein complexes between the training and test sets were excluded, resulting in 4,852 unique training complexes.

Data Augmentation

To increase the effective size of the training set, each protein-ligand complex was randomly rotated nine times using random rotation matrices, yielding ten variants per complex. All variants share the same PLA score, resulting in a total of 48,520 training samples.

Featurization

Protein-ligand complexes are represented as 3D grids of size $20 \times 20 \times 20$, with each grid cell encoded as a one-hot vector of length 28. This vector comprises 14 protein atom types² and 14 ligand atom types. The resulting training tensor has shape (48520, 20, 20, 20, 28).




Fig. 1. Featurization of protein-ligand complexes. Example shown: PDB ID 1a30. Default resolution is $20 \times 20 \times 20$ with 28 atomic categories.

Network Architecture

The original Sfcnn publication describes four network architectures and three featurization strategies. The architecture depicted in Figure 2, combined with the aforementioned featurization, achieved optimal validation performance.

¹ <https://bmcbioinformatics.biomedcentral.com/articles/10.1186/s12859-022-04762-3>

² Visit the original Sfcnn publication for atom type definitions

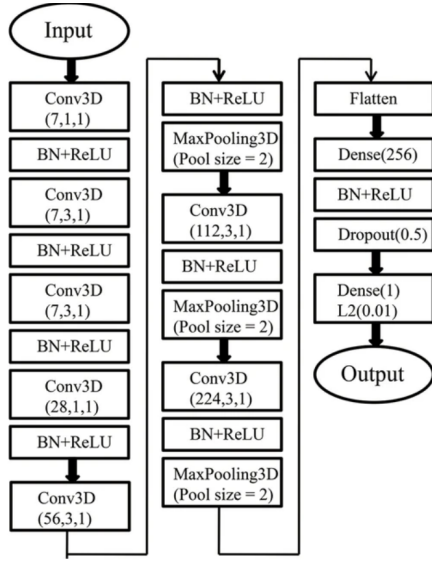


Fig. 2. Final CNN architecture for the Sfcnn network.

This architecture employs 3D convolutional layers with batch normalization and ReLU activation. L2 regularization is applied to the output layer to mitigate overfitting and enhance generalization.

Reproduction Study

Dataset and Featurization

The reproduction pipeline utilizes the same datasets and featurization approach, yielding a training tensor of shape (48520, 20, 20, 20, 28) and a test tensor of shape (285, 20, 20, 20, 28).

Data Storage

The original Sfcnn implementation stored data as concatenated arrays in a single .pkl (pickle) file, requiring all data to reside in memory, which is impractical for extremely large datasets. We adopted the .h5 (HDF5) format via h5py, enabling incremental writing and efficient storage. The resulting training grid occupies 40.1 GiB.

Network Implementation

The PyTorch implementation closely mirrors the original TensorFlow version, with two key differences:

1. Due to PyTorch’s Conv3D API, input tensors are permuted to shape (batch_size, 28, 20, 20, 20).
2. PyTorch lacks a direct L2 regularization API; instead, weight decay is applied to the final fully connected layer to approximate this effect.

Training Procedure

Training and validation sets are partitioned as in the original study, with the validation set comprising indices 41,000 to 48,520. The resulting splits are: training (41,000, 20, 20, 20, 28), validation (7,520, 20, 20, 20, 28), and test (285, 20, 20, 20, 28), corresponding to a ratio of 84.00% : 15.42% : 0.58%. The original hyperparameters did not yield convergence

in our PyTorch experiments; both original and adjusted hyperparameters are summarized in Table 1.

Table 1. Original and Reproduced Hyperparameters

Parameter	Original	Reproduced
Learning rate	0.004	0.0015
Batch size	64	32
Dropout rate	0.5	0.15
L2 regularization / FC weight decay	0.01	0.01
Epochs	200	200

Results

Evaluation Metrics

Sfcnn performance is evaluated using the following metrics:

$$\begin{aligned}
 \text{RMSE} &= \sqrt{\frac{1}{N} \sum_{i=1}^N (y_{\text{predict}} - y_{\text{true}})^2} \\
 \text{MAE} &= \frac{1}{N} \sum_{i=1}^N |y_{\text{predict}} - y_{\text{true}}| \\
 \text{SD} &= \sqrt{\frac{1}{N-1} \sum_{i=1}^N ((ay_{\text{predict}} + b) - y_{\text{true}})^2} \\
 R &= \frac{\mathbb{E}[(y_{\text{predict}} - \mu_{y_{\text{predict}}})(y_{\text{true}} - \mu_{y_{\text{true}}})]}{\sigma_{y_{\text{predict}}} \sigma_{y_{\text{true}}}}
 \end{aligned}$$

where a and b are the slope and intercept of the linear regression between predicted and measured values, $\mathbb{E}[\cdot]$ denotes expectation, and μ and σ represent means and standard deviations, respectively.

Table 2. Performance Metrics on CASF-2016 Core Set

Metric	Reproduced Sfcnn	Original Sfcnn
Pearson R	0.7286	0.7928
RMSE	1.5481	1.3263
MAE	1.2579	1.0277
SD	1.4892	1.3252

Although the original Sfcnn reports superior metrics, its training process did not converge in our reproduction, raising concerns about its reported performance. Due to the lack of access to the original training data and the absence of author response to data requests on GitHub³, the original training process is deemed irreproducible.

The training, validation and testing result on the four metrics are shown in the figure below.

³ <https://github.com/bioinfocqupt/Sfcnn/issues/1>

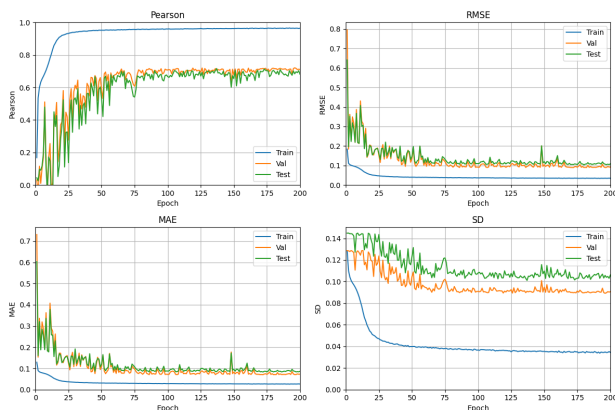


Fig. 3. Training curve for reproduced hyperparameters.

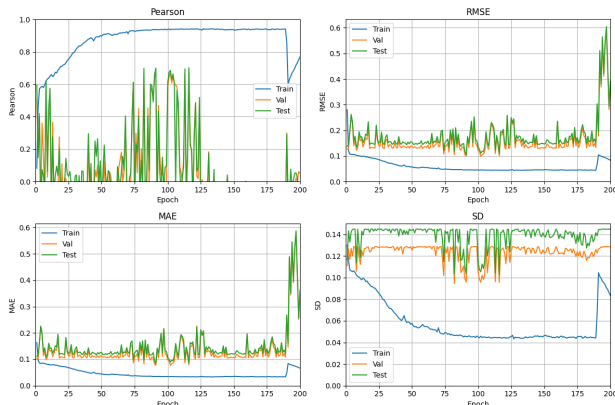


Fig. 4. Training curve for original hyperparameters.

The convergent results from our reproduction are considered the baseline for subsequent AF3 assessment.

AlphaFold3 Structure Assessment

Dataset

The assessment utilizes the CASF-2016 core set, excluding six proteins with structural complexity beyond AlphaFold3's predictive capacity, resulting in 279 proteins. Proteins with more than five isomorph or heterogeneous chains were excluded, as detailed in Table 3.

Table 3. Excluded Complex Protein Structures

PDB ID	Number of Chains
2xb8	12
2ymd	10
3n76	12
3n7a	12
3n86	12
4ciw	12

Workflow

Structure Generation

Protein structures were generated using the *Chai-1 online server*⁴. The AlphaFold3 online server⁵ was not used due to its inability to accept specific ligand SMILES codes. The MSA (Multiple Sequence Alignment) option was enabled with the MMseqs2 algorithm for all generations.

Docking and File Handling

Server outputs were downloaded as zip archives containing multiple ranked structures and associated metrics. The top-ranked structure (**pred.rank.0.cif**) was selected for analysis. To avoid conversion errors between file formats (.cif, .pdb, .mol2), structures were parsed directly using the **MMCIFParser** from Bio.PDB, followed by featurization and grid mapping. Note: Atoms or isotopes not included in the 14 predefined atom types are categorized as **other**.

Scoring

Predicted structures were scored using the reproduced Sfenn network, loaded with pre-trained weights (Pearson R = 0.728). Detailed analysis of PLA results and metrics is provided in the subsequent section.

Conclusion

This section will summarize the principal findings and implications of the study. [Please insert your conclusions here.]

External Libraries

- **PyTorch**: Custom neural network implementation.
- **Pandas**: Data storage and analysis.
- **Numpy**: Data processing and manipulation.
- **Matplotlib**: Visualization of training loss curves.

Author Contributions

G.Y., L.T., and J.C. conceived, conducted, and analyzed the experiments, and contributed to manuscript preparation and review.

Acknowledgments

The authors thank the anonymous reviewers for their valuable feedback. This work was supported in part by the National Science Foundation (NSF: #1636933 and #1920920).

⁴ <https://lab.chaidiscovery.com/dashboard>

⁵ <https://alphafoldserver.com/>